

Introduction to Data Science
CSC59970 | CCNY
Professor Grant M. Long

Final Project
project_finding.md

Abdur Rafey
Dzhonibek Parmankulov
Hasibul Islam

Questions and Tasks:

(1) Data Usage.

(a) What outside data have you appended to the original data set? Why did you choose this data?

- We decided to append apartment zip codes and income. We wanted to use the zip codes to see if the income of tenants have anything to do with the rents at the areas.

(b) Does the inclusion of this additional data raise any ethical considerations?

- One ethical consideration could be that the zip codes act like a sector divider. Thus, we get an idea of the common rent prices in different areas of New York City, which gives an indication of the social classes of tenants living in those areas.

(2) Data Exploration.

(a) What outliers present issues for your analysis? How have you chosen to handle them? Why?

- Outliers are luxury apartments with very high rent price and cheaply rented apartments.

(b) To what extent do missing values pose a challenge for your analysis? How have you chosen to handle them? Why?

- Missing values didn't really affect our results too much, but they do change the outlook of the conclusion of our analysis if not handled. We decided to use the median values of the other available data for the specific columns that had null values and filled them with those medians. We choose median over other statistics like mean because possible outliers don't have much effect at all on the median, whereas they can skew the mean quite a bit and possibly give a false analytical conclusion.

(c) Are there any other aspects of the data your exploration shows might be problematic?

- The problem with our model is that, due to the missing data in the training set, we were forced to use medians of the available data to fill in the null values so our final results don't completely reflect what a perfect model would be like, but it is still fairly close to the actual results.

(d) Create at least one visualization that demonstrates the predictive power of your data.

OLS Regression Results						
=====						
Dep. Variable:	rent	R-squared (uncentered):	0.836			
Model:	OLS	Adj. R-squared (uncentered):	0.836			
Method:	Least Squares	F-statistic:	6487.			
Date:	Thu, 05 Dec 2019	Prob (F-statistic):	0.00			
Time:	22:21:22	Log-Likelihood:	-1.2528e+05			
No. Observations:	14000	AIC:	2.506e+05			
Df Residuals:	13989	BIC:	2.507e+05			
Df Model:	11					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

bedrooms	145.6283	19.185	7.591	0.000	108.024	183.233
year_built	1.4243	0.091	15.662	0.000	1.246	1.603
bathrooms	1622.6973	36.597	44.340	0.000	1550.963	1694.432
min_to_subway	-0.0023	0.006	-0.355	0.722	-0.015	0.010
size_sqft	2.4718	0.040	61.379	0.000	2.393	2.551
no_fee	-232.1212	34.133	-6.801	0.000	-299.026	-165.216
has_doorman	569.6259	53.495	10.648	0.000	464.768	674.483
addr_zip	-0.3815	0.016	-23.427	0.000	-0.413	-0.350
floor_count	18.6267	1.911	9.749	0.000	14.882	22.372
has_gym	498.2204	51.740	9.629	0.000	396.802	599.638
allows_pets	297.7480	34.079	8.737	0.000	230.950	364.547
=====						
Omnibus:	13693.453	Durbin-Watson:	2.023			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	3379523.106			
Skew:	4.200	Prob(JB):	0.00			
Kurtosis:	78.650	Cond. No.	4.32e+04			

(3) Transformation and Modeling.

(a) Describe 5-10 features you think play the biggest role in your model.

These were the features that played the biggest role in our model

- bedrooms
- bathrooms
- size_sqft
- addr_zip
- floor_count

Number of bedrooms/bathrooms and floor_count had a positive correlation to the rent. Also, the addr_zip is a feature we added to divide our data into sectors so we could see what incomes are most prevalent in which areas.

• How did you create these features?

- We created these features by importing the training dataset and for the zip codes, we used an external dataset from data.cityofnewyork.us.

• How do you know these features are playing key roles?

- These features are playing key roles because as bedrooms or bathrooms increase, the rent goes up, especially in certain zip codes. Another feature to take into account is best nearest restaurant. We used grades of the best nearest restaurant to see how rent increases/decreases based on this feature.

(b) Describe how you are implementing your model. Why do you think this works well?

- WE are implementing it using features we chose and predicting the target which is rent. Initial training dataset with 12000 rental records and then we predicted 2000 records of test data to check for accuracy/mean squared error.

(c) Describe your methodology for selecting your model. Why do you think this type of model works well?

- Initially we started by using Linear Regression because we knew that we had continuous data to work with so that seemed to be the best option. Later, we tried Forest Tree and it gave a much smaller Mean Squared Error result due to its ability to navigate through more complex and larger datasets. With this idea in mind, we decided to use gradient boosting and it worked in a very similar method and gave a slightly better result.

(4) Metrics, Validation, and Evaluation.

(a) How well do you think your model will perform on the hold out test set? How do you know? 3

- Linear Regression: Mean Squared Error using Linear Regression: 3313817.143868871
Explained variance regression score function Best possible score is 1.0, lower values are worse.
0.5559738486087318
- R² (coefficient of determination) regression score function: 0.5556196902015677

Random Forest: Mean Squared Error using Random Forest Regressor: 1883630.165411068
Explained variance regression score function Best possible score is 1.0, lower values are worse.
0.7474068681056247

- R^2 (coefficient of determination) regression score function: 0.7474066551922682

Gradient Boosting: Mean Squared Error using GradientBoostingRegressor: 1720228.7848562498
Explained variance regression score function Best possible score is 1.0, lower values are worse.
0.7693223390436726

- R^2 (coefficient of determination) regression score function: 0.7693186536399759

We think our model would perform fairly well on the hold out test set. Based on the results for each model we tried, we got closer to a perfect score of 1, so our model was pretty consistent and our mean squared error decreased with each trial and gradient boosting gave the best results.

(b) Is your model useful? Why or why not?

- Our model is useful as its variables correlate with each other at a fairly consistent rate based on the R-squared value.

(c) Are there any special cases in which your model works particularly well or particularly poorly?

- Our model works well on rentals with average rent price, but it is poorer for luxury rentals.

(d) Create at least one visualization that demonstrates the predictive power of your model.

OLS Regression Results						
=====						
Dep. Variable:	rent	R-squared (uncentered):	0.836			
Model:	OLS	Adj. R-squared (uncentered):	0.836			
Method:	Least Squares	F-statistic:	6487.			
Date:	Thu, 05 Dec 2019	Prob (F-statistic):	0.00			
Time:	22:21:22	Log-Likelihood:	-1.2528e+05			
No. Observations:	14000	AIC:	2.506e+05			
Df Residuals:	13989	BIC:	2.507e+05			
Df Model:	11					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

bedrooms	145.6283	19.185	7.591	0.000	108.024	183.233
year_built	1.4243	0.091	15.662	0.000	1.246	1.603
bathrooms	1622.6973	36.597	44.340	0.000	1550.963	1694.432
min_to_subway	-0.0023	0.006	-0.355	0.722	-0.015	0.010
size_sqft	2.4718	0.040	61.379	0.000	2.393	2.551
no_fee	-232.1212	34.133	-6.801	0.000	-299.026	-165.216
has_doorman	569.6259	53.495	10.648	0.000	464.768	674.483
addr_zip	-0.3815	0.016	-23.427	0.000	-0.413	-0.350
floor_count	18.6267	1.911	9.749	0.000	14.882	22.372
has_gym	498.2204	51.740	9.629	0.000	396.802	599.638
allows_pets	297.7480	34.079	8.737	0.000	230.950	364.547
=====						
Omnibus:	13693.453	Durbin-Watson:	2.023			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	3379523.106			
Skew:	4.200	Prob(JB):	0.00			
Kurtosis:	78.650	Cond. No.	4.32e+04			
=====						

(5) Conclusion

(a) How would you use this model?

- We can use this model to determine if one sector of NYC would be best suited to your needs depending on the contributing environmental factors of that apartment location. For example, a higher rent based on better restaurants or close train stations can be seen differently from higher rent based on number of bedrooms/bathrooms.

(b) If you could have additional modeling features, what would they be?

- Since we are mostly taking into account environmental factors as the additional data for our experiment, we could also choose to use number of playgrounds near apartment or best rated schools near apartment if the tenants have children.

(c) Would you rather have more data, or more features?

- We would rather have more data because sometimes more data can give us an opportunity to test out more features and see the differences in the results if any. Also, having more data can also bring the accuracy closer to the expected value so using more features after adding extra data into a model could produce better results. This is similar to the law of large numbers. As you perform more tests you get closer to an expected value.